PREDICTIVE CODING OF GENOMIC INFORMATION DATA

BY

Jan Voges
Erichstr. 2
30449 Hannover
Germany
Citizenship: German

Marco Munderloh
Musterstr. 42
3XXXX Musterstadt
Germany
Citizenship: German

Joern Ostermann
Karla-Schmidt-Str. 14
30655 Hannover
Germany
Citizenship: German

Contact
{jvoges, munderl, office}@tnt.uni-hannover.de

Table of Contents

1   Preface

It should be understood at the outset that although an illustrative implementation of one or more embodiments are provided below, the disclosed systems and/or methods may be implemented using any number of techniques, whether currently known or in existence. The disclosure should in no way be limited to the illustrative implementations, drawings, and techniques illustrated below, including the exemplary designs and implementations illustrated and described herein, but may be modified within the scope of the appended claims along with their full scope of equivalents.

2   Introduction

Due to novel high-throughput next-generation sequencing (NGS) technologies, the sequencing of huge amounts of genetic information has become affordable. On account of this flood of data, IT costs may become a major obstacle compared to sequencing costs. High-performance compression of genomic data is required to reduce the storage size and transmission costs.

Raw sequencing data (mainly the base sequences – also known as reads - and the quality scores) is stored in so-called FASTQ files, whereas mapped sequence data is stored in so-called SAM files. SAM files are plain-text human-readable files containing the reads and quality scores as well as mapping information with respect to some reference genome. Mapped sequencing data represented in SAM files contains more redundancy, as typically multiple reads are mapped to the same location on the reference genome. The average amount of reads mapping to the same location is referred to as *coverage*.

It has been shown by Tembe et al. [7] and Deorowicz and Grabowski [8], that splitting the data into separate streams for sequence reads, quality scores etc. (and compressing them independently) yields significant gains over general-purpose (dictionary-based [2]) programs like *gzip*. All recent contributions to the topic of FASTQ and/or SAM file compression employ this scheme of separately compressing the sequence reads and quality scores (as well as the other field present in FASTQ or SAM files).

Based on statistical analysis performed on a reference data set (issued by the appointed MPEG ad-hoc group on genome compression), new proposals for sequence and quality score compression in SAM files have been developed.

2.1   Outline

We propose a sequence compressor which assumes aligned and position-sorted data. Our sequence compressor combines the mapping positions, the CIGAR strings and the actual nucleotide sequences to implicitly assemble local parts of the donor genome and compress the sequence reads. The compressor does not need a reference genome and is able to perform compression "on-the-fly" using solely a *sliding window* as context for the local assembly.

Regarding the compression of quality scores, we propose some compression schemes based on Markov models with a variable context to boost prediction performance. Due to the large alphabet of quality scores (and thus a huge memory consumption of any Markov model), we furthermore propose employing FIR filter-based compression for quality scores.

3 Block-based Compression of Nucleotide Sequences

3.1 Prior Work

Current implementations (e.g. *Quip* [6] and *sam_comp* [3]) use an order-2 arithmetic coder (AC) to compress the nucleotide sequences. In order to exploit the redundancy present in the data, an arithmetic coder needs *a priori* knowledge. This practically excludes random access to the compressed data as the prediction performance increases with the amount of data the Markov model can be trained on.

The authors of *DeeZ* [5] made a new approach, called "local assembly" (implicit assembly of the donor genome using the mapping information present in SAM files). *DeeZ* nevertheless performs the local assembly with respect to an external reference (given in FASTA format).

The authors of *Quip* [6] implemented an assembly-based compression scheme for nucleotide sequences, too. They use by default the first 2.5 million reads to assemble so-called contigs which are then used in place of a reference sequence to encode aligned reads. This algorithm does not require a reference – this means that the resulting compressed files are entirely self-contained – but has the drawback that the best match of a specific read to the previously assembled reference has to be found.

3.2 Proposed Compression Mechanism

We propose a compression scheme using a sliding window to perform a "local block-based implicit assembly". In contrast to *DeeZ*, our compressor does not require a reference genome and is able to work on smaller block granularities as it uses a "sliding window" as context (of course smaller block sizes have a negative impact on compression ratio). In contrast to *Quip*, we do not need any reads to assemble so-called contigs in advance. We perform the compression using only the currently available nucleotide sequences in the sliding window. Our compression scheme therefore is less complex but should nevertheless yield comparable compression ratios.

Our approach is based on correlations between consecutive reads. As all reads in a SAM file are aligned to their mapping position, all reads mapped to the same position should be similar, except for gene mutations (SNP's) or sequencing errors respectively mapping errors. These aberrations are coded in the CIGAR string. An example is given in figure 1 and figure 2. Consecutive reads are plotted line by line. All reads mapped to the same position are aligned. As SAM files are stored sorted by mapping position, a read mapping to a different location than the read before leads to a step in the plot. If plotted over a large block, this leads to the staircase property visible in the plotted figure(s).

Figure 1: Consecutive nucleotide sequences from a SAM file. The different colors represent the nucleotides.


Figure 2: The nucleotide sequences after the `expand` function, stored as a sliding window in the circular buffer. The noisy 'dots' represent genome mutations or sequencing errors described in the CIGAR string.

Our approach proposes to encode sequence reads block-wise exploiting the joint coding of the mapping positions, the CIGAR strings and the sequence reads itself, whereas the block size might be fixed or variable and of arbitrary size. The first sequence read in a block is encoded without prediction, as well as the corresponding mapping position and the CIGAR string. However, it might also be coded against some reference. Some subsequent read $i, i > 1$, is then aligned to a previous read $j, j < i$, using its CIGAR string and its mapping position. Thereafter, we compute the difference d (containing the differences of the mapping position and the nucleotide sequence from the aligned read i to read j) and pass it to an entropy coder, e.g. an order-0 arithmetic coder. At the decoder, read i can be reconstructed from read j by applying d to reconstruct the nucleotide sequence, the CIGAR string, and the mapping position.

The alignment of read i to read j is done with a function called `expand`. The CIGAR strings and mapping positions of both reads are used to unwind the sequence such that they would both directly align to some external reference (please refer to the SAM file specification [1] for a detailed explanation of the CIGAR string syntax). The output of the `expand` function is given in Figure 2. Consequently, `expand` condenses the position, the CIGAR string, and the sequence into a joint representative code word. As an alternative, these SAM fields could be compressed using separate entropy coders.

A short example of the working principle of `expand` is given in Figure 3.

```
POS CIGAR       SEQ
7    8M2I4M1D3M TTAGATAAAGGATACTG < read 1
9    3S6m1P1I4M AAAAGATAAGGATA    < read 2


     6789... < POS
      8M       2I4M  1D3M
      TTAGATAAAGGATA_CTG < read 1 expanded
     AAAAGATAAGGGATA     < read 2 expanded
     3S 6M      1P1I4M
```
Figure 3: Working principle of the function `expand`

The previous read j is by default read $i - 1$. This context read $j = i - 1$ might be highly erroneous or not aligned to the reference. We address this issue by keeping track of N previous reads (sliding window) to be able to select the best matching read to encode read i. Sequence reads which are not aligned to the reference might be directly passed to the entropy coder or retained and encoded

separately later, e.g. at the end of each block. The outline of the proposed algorithm is shown in Figure 4.

```
Input: block_size, threshold, pos[L], cigar[L], seq[L], L (total    no. of SAM
records)
Output: binary stream (from entropy coder)
new circbuf(N); // circular buffer of size N (sliding window)

l = 0;         // line count
b = block_size; // block line count

while (l < L)
| if (b == block_size) then // start new block
| | clear_circbuf();
| | reset_entropy_coder();
| | entropy_coder(pos[l], cigar[l], seq[l]);
| | circbuf.push(pos[l], cigar[l], seq[l]);
| | b = 1;

| else
| | if (cigar(l) != '*') then // sequence is mapped
| | | min_entropy = +inf;
| | | do
| | | | seq_exp = expand(pos[l], cigar[l], seq[l]);
| | | | (pos_ref, cigar_ref, seq_ref, ref_index) = circbuf.get_next();
| | | | ref_exp = expand(pos_ref, cigar_ref, seq_ref);
| | | | d = compute_differences(seq_exp, ref_exp);
| | | | if (entropy(d) < min_entropy) then

| | | | | min_entropy = entropy(d);
| | | | | ref_index_best = ref_index;
| | | | | d_best = d;
| | | | endif
| | | while (entropy(d) > threshold) && circbuf.has_next();
| | | entropy_coder(ref_index_best, d_best);
| | else
| | | entropy_coder(pos[l], cigar[l], seq[l]);
| | endif
| | circbuf.push(pos(l), cigar(l), seq(l));
| endif
| b++;
| l++;
endwhile
```

Figure 4: Algorithm outline

Note: Per block, the initial read and the trailing substrings of subsequent reads can be regarded as locally assembled genome.

4 Predictive Coding of Quality Scores

4.1 Prior Work

The quality scores yield the largest first-order entropy among all field present in a typical SAM file, mainly due to their large alphabet. The authors of *Fqzcomp* [3] state the following dependencies for a sequence of quality scores:

- Any score q_i has a strong correlation to the direct previous quality score $q_{i-1}, q_{i-2}, \dots$, decreasing the further back we go [4].
- A known issue with the Illumina base-caller causes many sequences to end with quality score 2 ("#").
- There is a correlation between position and quality values. Quality values are typically reducing along the length of the sequence.
- Sequences as a whole tend to be good or bad.

The authors of *Fastqz* [3] additionally state that it was observed that the quality values tend to start with a common maximum value.

Current SAM compression implementations exploit this observations by employing a Markov model with the memory ranging over the direct predecessors in the quality score stream.

4.2   Modeling the Source with a Markov Chain

In terms of performing some statistical analysis regarding the quality scores, we converted the quality scores into single streams. The auto-correlation functions of these streams do not resemble a delta impulse and thus the power spectral density is not constant. This means, that the quality score source has some kind of memory. Thus, it is highly suitable to model the source with a Markov chain.

Common SAM file compression approaches [4,5] noted that quality scores usually contain a trailing sequence of low quality values. We made this observations, too, and propose first to trim the trailing low quality values.

Current SAM file compressors model the quality score source as a Markov chain, predicting the current symbol with a maximum likelihood predictor using the *N* direct predecessors.

The maximum likelihood predictor selects in each state $q_1, \ldots q_N$ the symbol a_i with the greatest conditional probability *P* as prediction value $q_{N+1,p}$:

$$P\big(q_{N+1,p} = a_i \big| q_1, \ldots q_N\big) \geq P(q_{N+1} = a_k | q_1, \ldots, q_N) \forall k$$

4.2.1   Drawbacks and Recommendations

The maximum likelihood predictor consumes a huge amount of memory since we have to keep track of $k^N * k$ relative symbol frequencies (having an alphabet of size k). This model also needs a large amount of data to train in order to tightly fit the model to the data. We propose to store the prediction table (generated during compression) in the file header to speed up decompression. Prediction tables for the individual blocks could then be delta-encoded. This would enable random access to the compressed data while retaining the compression performance of the Markov model.

4.3   Proposed Compression Schemes

4.3.1   Extended Line Context Approach

The authors of *Fqzcomp* [3] state, that quality score lines as a whole tend to be "bad" or "good". Additionally, we made the observation, that in some datasets subsequent quality score lines are "similar" in terms of a small Levensthein distance between pairs of quality score lines.

We respond to this using a so-called line context, i.e. a structure holding some number of previous lines from the current block (this is also known as "sliding window"). The Markov memory (i.e the prediction context) is divided into one part containing the memory locations on the current line ("intra-line memory") and one part containing the memory locations on some specific line from the line context ("inter-line memory").

Choosing the right line from the line context to boost prediction performance is crucial. We compute a three-dimensional identification vector for each quality score line, containing its length, its arithmetic mean and its variance. In the process of encoding line *i*, we first select the one specific line $\hat{\imath}$ from the line context with the least identification vector distance. Consecutively, the intra-line memory is placed on line *i* and the inter-line memory is placed on the selected line $\hat{\imath}$ from the line context.

This compression approach yields a relatively low encoder complexity yet employing contexts ranging over multiple quality score lines.

In the following sections, we propose two additional approaches to exploit the line context.

4.3.2 Dictionary-based Substring Matching Approach

To encode the quality score $q_{i,j}$ in column j and line i, we propose to select the substring s of N preceding symbols $s = q_{i,j-N}, \dots, q_{i,j-1}$. We then propose to search for a similar substring $\hat{s} = q_{\lambda,\gamma}, \dots, q_{\lambda,\gamma+N-1} \equiv s$ in the line context and to use $q_{\lambda,\gamma+n}$ as the prediction value $q_{i,j}^{(p)}$ for $q_{i,j}$.

Furthermore, we propose to use a prediction value $q_{i,j}^{(p)}$ calculated as a mean or median from M substring matches $\hat{s}_m$, which might be weighted by their prediction error $\left| q_{i,j}^{(m)} - q_{i,j}^{(p,m)} \right|$, whereas any suitable norm might be applied.

4.3.3 Markov Context-Extending Substring Matching Approach

As an alternative, we propose to extend the context of any of the above mentioned Markov predictors with the prediction value $q_{\lambda,\gamma+N}$ found by substring matching within the line context.

4.4 Linear Prediction with FIR Filters

Due to the large alphabet of the quality scores (typically ~30 symbols, although more are theoretically possible), Markov models are limited to some low order N (typically $N \leq 3$).

As an alternative, we propose to use an FIR filter to predict the current symbol. While it is on the one hand costly to compute the FIR filter coefficients, on the other hand, filter coefficients might be shared between multiple quality score lines.

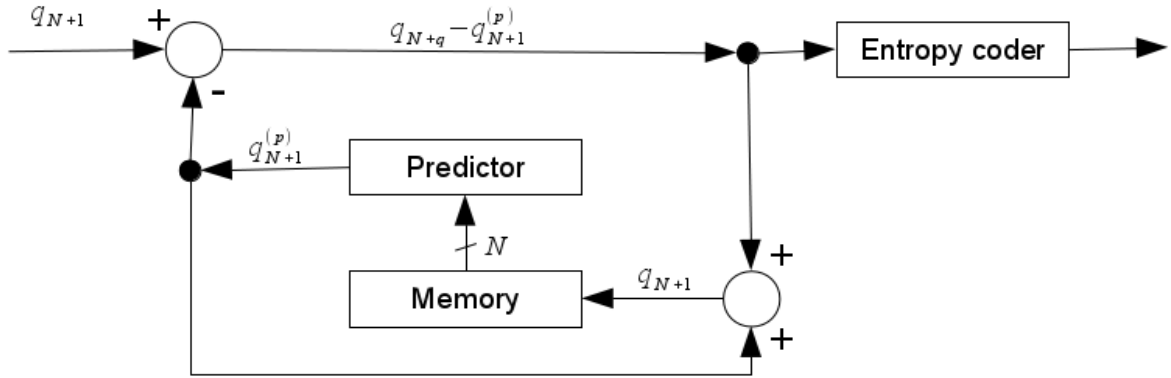The basic layout of a predictive encoder can be viewed in Figure 5.

Figure 5: Linear predictive encoder schematic

The linear predictor computes the current prediction value $q_{N+1}^{(p)}$ by evaluating the linear series of previous symbols q_n weighted with predictor coefficients a_n.

$$q_{N+1}^{(p)} = \sum_{n=1}^{N} a_n q_n$$

The usual approach to compute the predictor coefficients is to minimize the mean squared error which yields the following condition to compute the filter coefficients.

$$\frac{\partial}{\partial a_n}\left(E\left[\left(u_{N+1}-u(N+1)^{(p)}\right)^2\right]\right)=0$$

Thus, we obtain l sets of predictor coefficients for a quality score line of length l. Using the line context from section 3.3, it could be beneficial to share previously computed predictor coefficients among multiple lines.

In contrast to this minimum-mean-square-error filtering it might be beneficial to investigate and employ minimum-variance prediction and filtering.

4.5   Entropy Coding of Prediction Errors

The histogram of the prediction errors of the Markov or FIR predictor approximately shows a two-sided geometric distribution. We therefore propose to use subsequent Rice coding [9] as an alternative to arithmetic coding to obtain the binary representation of the data. This reduces the encoder complexity, as recent compressors such as *Quip* [6] employ arithmetic coding.

5   References

[1]   The SAM/BAM Format Specification Working Group: Sequence Alignment/Map Format Specification.

[2]   Ziv., J. & Lempel, A.: A universal algorithm for sequential data compression. IEEE Transactions on information theory 23, 337-343 (1977).

[3]   Bonfield, J.K & Mahoney, M. V.: Compression of FASTQ and SAM Format Sequencing Data. PloS ONE 8, e59190 (2013).

[4]   Christos Kozanitis, Chris Saunders, Semyon Kruglayak, Vineet Bafna, and George Varghese: Compressing genomic sequence fragments using SlimGene. Journal of Computational Biology: A Journal of Computational Molecular Cell Biology, 18(3):401-13, March 2011. ISSN 1557-8666. doi: 10.1089/cmb.2010.0253.

[5]   Faraz Hach, Ibrahim Numanagic & S. Cenk Sahinalp: DeeZ: reference-based compression by local assembly. Nature Methods, Vol. 11, No. 11, November 2014, pp 1082-1084.

[6]   Daniel C. Jones, Walter L. Ruzzo, Xinxia Peng & Michael G. Katze: Compression of next-generation sequencing reads aided by highly efficient de-novo assembly. Nucleic Acids Research, 2012, 1-9. doi: 10.1093/nar/gks/754.

[7]   Tembe, W., Lowey, J. and Suh, E.: G-SQZ: compact encoding of genomic sequence and quality data. Bioinformatics, 26, 2192-2194, (2010).

[8]   Deorowicz, S. and Grabowski, S.: Compression of DNA sequence reads in FASTQ format. Bioinformatics, 27, 860-862. (2011).

[9]   Robert F. Rice and James R. Plaunt: Adaptive Variable-Length Coding for Efficient Compression of Spacecraft Television Data, IEEE Transactions on communication technology, Vol. COM-19, No. 6, December 1971.